


Predictions were saved before the outcome

Evaluation model: gpt-oss-120b.

The study measured performance on authored coding tasks.



Hashes support checking for changes to saved files.

Local files can still be edited outside the workflow.

No human learning outcome was measured.