

Which parts of the project run where?

The local demo is deterministic in the path shown here.

The benchmark calls the model. Analysis reads its saved responses.

Interactive proof-of-concept demo

One local tutoring turn.

Person using demo

Sends an answer to a prepared task.



React + FastAPI

Shows the session and checks requests.



LangGraph turn

Updates the tracker and chooses an action.



Saved turn log

Saves each turn for replay.



Prompt template

Uses a template to ask the next question.

Sealed external-model benchmark

Predictions saved before outcome request.

Prepared cases

Separates three task views.



Evaluation model

Supplies evaluation responses.



Code check + tracker

Runs probe code and records results.



Outcome code check

Runs outcome code and records completion.



Outcome request

Requests the later task after predictions.



Saved predictions

Saves four predictions before outcome.

Offline analysis and tracker study

Analyses saved records without external calls.

Saved benchmark

Stores predictions and reviews.



Result comparison

Compares paired results.



Report figures

Creates report files.

Controlled simulator

Creates states and noisy evidence.



Five Bayesian trackers

Compares update rules.



Tracking results

Reports prediction error and runtime.

Results concern prepared cases and simulated states.

Neither study measures human learning.