


The corrected probe changed one decision and required more work

The recorded responses were checked again after repairing a test-harness error; no model or sandbox call was repeated.


The relevant probe changed one of 23 eligible decisions from wrong to right.

The probe condition used one extra model response and one code execution for each case.

A. Corrected predictions



B. Work used for each prediction



C. Work observed in the full run

72
model responses

32,583
input + output tokens

32.7 s
summed request latency

47/48
completed sandbox attempts

2.27 MB
measured artefacts

Not reported
provider cost

Post-hoc result: the correction shows that the original result depended on the test harness; it does not replace the sealed analysis.

Scope: one pinned evaluation model, one run per authored case. Provider cost and sandbox CPU, memory, and execution time were not recorded.

These measurements describe this experiment's workload; they do not show deployment cost, learning, or tutoring effectiveness.