


A test-harness error changed the benchmark result

Recorded responses were rescored after repairing the tests.
No new model responses or sandbox executions were collected.


The repair changed 19 of 48 execution records.

After correction, adding evidence gave one extra correct prediction.


A. Prediction accuracy changed



B. Estimated advantage fell



C. Comparison with control evidence



The interval includes no improvement. Related and unrelated passing checks produced identical predictions.

One model run on 23 eligible authored cases.
This does not show human learning or tutoring effectiveness.